

# **Weighted Priority Arbiter Implementation**

**March 2014**

© 2004-2014 Cadence Design Systems, Inc. All rights reserved.  
Printed in the United States of America.

Cadence Design Systems, Inc. (Cadence), 2655 Seely Ave., San Jose, CA 95134, USA.

Product JasperGold Apps incorporates software developed by others and redistributed according to license agreement. For further details, see doc/third\_party\_readme.txt.

**Trademarks:** Trademarks and service marks of Cadence Design Systems, Inc. contained in this document are attributed to Cadence with the appropriate symbol. For queries regarding Cadence's trademarks, contact the corporate legal department at the address shown above or call 800.862.4522. All other trademarks are the property of their respective holders.

**Restricted Permission:** This publication is protected by copyright law and international treaties and contains trade secrets and proprietary information owned by Cadence. Unauthorized reproduction or distribution of this publication, or any portion of it, may result in civil and criminal penalties. Except as specified in this permission statement, this publication may not be copied, reproduced, modified, published, uploaded, posted, transmitted, or distributed in any way, without prior written permission from Cadence. Unless otherwise agreed to by Cadence in writing, this statement grants Cadence customers permission to print one (1) hard copy of this publication subject to the following conditions:

1. The publication may be used only in accordance with a written agreement between Cadence and its customer.
2. The publication may not be modified in any way.
3. Any authorized copy of the publication or portion thereof must include all original copyright, trademark, and other proprietary notices and this permission statement.
4. The information contained in this document cannot be used in the development of like products or software, whether for internal or external use, and shall not be used for the benefit of any other party, whether or not for consideration.

**Disclaimer:** Information in this publication is subject to change without notice and does not represent a commitment on the part of Cadence. Except as may be explicitly set forth in such agreement, Cadence does not make, and expressly disclaims, any representations or warranties as to the completeness, accuracy or usefulness of the information contained in this document. Cadence does not warrant that use of such information will not infringe any third party rights, nor does Cadence assume any liability for damages or costs of any kind that may result from use of such information.

**Restricted Rights:** Use, duplication, or disclosure by the Government is subject to restrictions as set forth in FAR52.227-14 and DFAR252.227-7013 et seq. or its successor

---

## Contents

---

<u>1</u>	
<a href="#">Preface</a>	5
<u>2</u>	
<a href="#">Overview</a>	7
<a href="#">Weighted-Priority Arbiter Architecture</a>	7
<a href="#">Weighted-Priority Arbiter Example</a>	7
<u>3</u>	
<a href="#">Interface</a>	11
<a href="#">Interface Description</a>	11
<a href="#">Steps for Customizing Your Interface Requirements</a>	11
<a href="#">Interface Names</a>	12
<a href="#">Examples</a>	12
<a href="#">Example 1. Grant Remains Active Until Transfer</a>	12
<a href="#">Example 2. Multiple Pending Requests on the Same Port</a>	13
<a href="#">Example 3. Latency Considerations for Request and Grant Signals</a>	14
<u>4</u>	
<a href="#">Testplan</a>	17
<u>5</u>	
<a href="#">Requirements</a>	21
<u>6</u>	
<a href="#">Strategy</a>	25
<a href="#">Verification Plan</a>	25

## Weighted Priority Arbiter Implementation

---

<a href="#">Restriction Definitions</a>	25
<a href="#">Verification Steps</a>	26
<a href="#">What Ports Should You Monitor as Part of Your Requirements?</a>	26
<a href="#">Speeding Up Your Proof</a>	28

## [7](#)

<a href="#">Coverage</a>	31
--------------------------	----

<a href="#">Appendix: Techniques and Strategies</a>	33
---	----

<a href="#">Taking Advantage of Design Symmetry</a>	34
<a href="#">Arbiter Example</a>	34
<a href="#">Muxitplexer Abstraction</a>	35
<a href="#">FIFO Abstraction</a>	37
<a href="#">Example</a>	38
<a href="#">Induction</a>	40
<a href="#">Example</a>	40
<a href="#">Counter Abstraction</a>	41
<a href="#">Counter Induction</a>	41
<a href="#">Counter Reduction</a>	42
<a href="#">Separating Decoding Logic</a>	43

---


## Introduction

---

Arbiters are a critical component in systems containing shared resources. For example, a design containing a memory controller, where multiple memory agent clients share a common memory bus, requires an arbitration scheme to prevent more than one agent from accessing the bus at the same time. Similarly, for a bridge design where more than one input port shares a common output port, arbitration is needed to prevent dropping or corrupting data transported through the bridge.

**Note:** If the arbitration scheme is in error, then service to a high-priority client might be ignored, which could result in lost data or even a system crash.

**Figure 1-1 Arbiters and Shared Resources**



There are many different arbitration schemes ranging from the unfair priority scheme to the fair round-robin scheme, or a combination of priority with fairness. Although arbiters are used in multiple applications, their basic requirements are generally straightforward.

This application note describes how to verify weighted priority arbiters.

# **Weighted Priority Arbiter Implementation**

## Introduction

---

---

## Overview

---

In this application note you will learn about a weighted priority arbiter implementation. This Overview is useful for understanding the various terms and phrases we use throughout the remainder of this application note.

This overview includes the following sections:

- [Weighted-Priority Arbiter Architecture](#)
- [Weighted-Priority Arbiter Example](#)

## Weighted-Priority Arbiter Architecture

In many cases, the bandwidth is not intended to be allocated evenly among ports. Some ports are intentionally provided more bandwidth than others. However, to prevent starvation of the lower bandwidth ports, strict priority is not used. Consequently, we give a “weight” to each port to indicate the relative expected bandwidth that at the same time gives all ports a chance to be granted the bus. There are many ways to accomplish weighting. One way is to give a larger time-slot to a particular port, and other ways include manipulating the round-robin priority pointers.

This arbiter belongs to a class where you must know the weighting algorithm. In this section we consider weighted priority that involves a well-defined hardware implementation (as opposed to a statistical software implementation that cannot be formally proven).

## Weighted-Priority Arbiter Example


A round-robin priority arbiter gives ports 0 through 3 twice as much bandwidth as ports 4 through 15. The implementation involves doubling the number of pointer locations for ports 0 through 3, as shown in [Figure 2-1](#) on page 8.

## Weighted Priority Arbiter Implementation

### Overview

---

**Figure 2-1 Round-Robin Arbiter Example**



As shown in [Figure 2-1](#) on page 8, Port0 through Port3 appear on both halves of the ports. Therefore in this example, to verify the weight of each port without using statistics, we need an exact way of figuring out whether the implementation of this algorithm is correct. Checking whether this algorithm gives Port0 statistically twice as much bandwidth as Port4 is not the objective of functional verification. Here we can count the maximum number of ports before reaching the port we are checking. In this case, the requirements are as follows:

- When Port0 through Port3 issue a request, they should get a grant before getting nine other transfers.
- When Port4 through Port15 issue a request, they should get a grant before getting 18 other transfers.


Another possible implementation is such that each port has a credit limit based on its “weight.” A fresh round of round-robin starts with all ports having equal access until ports start running out of their credit. All ports with credit remaining will continue in a round-robin manner until either all ports run out of credit or all the ports with credit remaining do not issue a request. After that, the round-robin starts again. This class of weighted round-robin can require many cycles just to complete one round. It is also difficult to write a requirement that keeps track of the credit remaining and adjust the arbitration accordingly. Fortunately, this class of arbiter can be easily represented by two distinct functions. Let us consider the following example:



## Weighted Priority Arbiter Implementation Overview

---

**Figure 2-2 Credit-Based Weighted Round-Robin**



The round-robin is arbitrated over all ports, but the ports that run out of “weights” or “credit” are excluded from the arbitration. Therefore, the first set of requirements reflects fairness arbitration over all ports with credit. We put a stopat on the credit generation to ignore the generation logic for this requirement. For the next requirement, we verify that if there is a grant, the weight will decrease; if all the ports run out of credit, or the ports that still have credit do not issue any request, the credit for all ports will be refreshed. For detailed verification information, refer to the application note “Dynamic Priority Arbiter Implementation.”

## **Weighted Priority Arbiter Implementation**

### Overview

---

---

# Interface

---

In this section you will learn about various interface characteristics you will need to consider when customizing our example requirements for your specific design. The examples in this section illustrate a few common arbiter interface requirements. Although your specific arbiter interface requirements might differ from the examples in this section, they are a starting point that you can modify to fit your specific needs.

This section includes the following topics:

- [Interface Description](#)
- [Steps for Customizing Your Interface Requirements](#)
- [Interface Names](#)
- [Examples](#)

## Interface Description

In addition to the weighted priority requirements previously described in the Overview (see [“Overview”](#) on page 7), often there are other temporal interface requirements for the relationship of the client's request and the arbiter's grant signals. For example:

- Is there a minimum time interval after a request when a grant must not occur due to latency considerations within the design?
- Does the interface support queuing multiple requests on the same port while waiting for a grant?

## Steps for Customizing Your Interface Requirements

Use the following steps to define your specific arbiter's interface requirement.

1. Generally describe the operational sequence of events for the interface.
2. Graphically illustrate the interface behavior.

## Weighted Priority Arbiter Implementation Interface

---

3. Using the operational description and graphical waveform, define the detailed temporal requirements for the interface.

### Interface Names

We use the following names for interface signals during our discussion and in the modeling requirements:

req [n]	Request issued by a client to the arbiter
gnt [n]	Grant issued by the arbiter to a client
xfer [n]	Start of a transfer issued by a client
p1	Specific input request port under test
p2	Specific input request port under test
req2gnt_latency	Required latency time between req and gnt

### Examples

Refer to the following interface examples:

- [Example 1. Grant Remains Active Until Transfer](#)
- [Example 2. Multiple Pending Requests on the Same Port](#)
- [Example 3. Latency Considerations for Request and Grant Signals](#)

#### Example 1. Grant Remains Active Until Transfer

**Step 1:** Generally describe the operational sequence of events for the interface.

- a. A request (Req) is asserted, indicating that the client is ready to transfer data on a bus.
- b. The request (Req) stays asserted until a grant is provided, after which the request is de-asserted.
- c. The arbiter eventually asserts a grant (Gnt).


## Weighted Priority Arbiter Implementation Interface

---

- d. The grant (Gnt) must remain asserted until up to four cycles or until a start of transfer occurs (Xfer).
- e. The grant (Gnt) issued to any client must remain de-asserted from the time that a transfer begins (Xfer) until the transfer ends.

**Step 2:** Graphically illustrate the interface behavior.

**Figure 3-1 Request/Grant/Transfer Relationships for Example 1**



**Step 3:** Using our operational description and graphical waveform, define the detailed temporal requirements for the interface.

- a. Request should be asserted until grant is asserted, after which request should be de-asserted.
- b. Transfer should not be asserted until grant is provided (constraint).
- c. Grant should be asserted until either four cycles have elapsed or after a transfer is asserted.
- d. Grant should not be asserted until the bus is idle (transfer is done).
- e. Only one grant can be issued at one time.
- f. Grant should be asserted if there is any request pending and the bus is idle.
- g. Grant should not be asserted if there is no request pending for any port.

### Example 2. Multiple Pending Requests on the Same Port

**Step 1:** Generally describe the operational sequence of events for the interface.

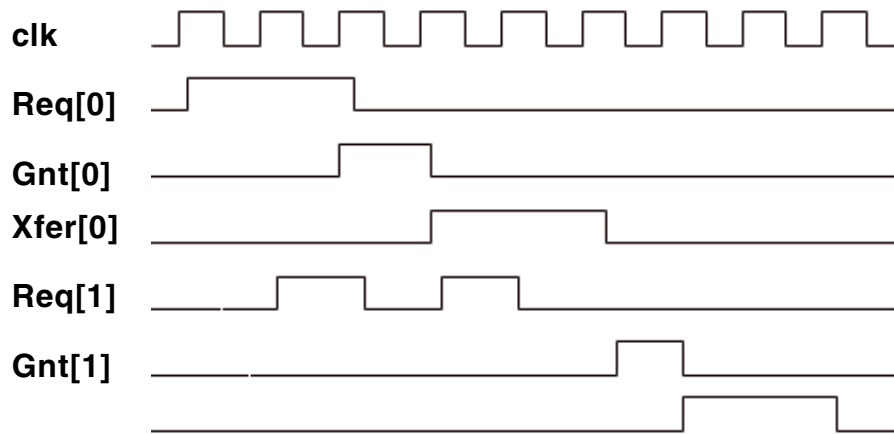
## Weighted Priority Arbiter Implementation Interface

---

- a. Request is asserted for one cycle, indicating that the client is ready to transfer data.
- b. Request can be asserted again. Up to four pending requests can be issued.
- c. Arbiter asserts grant for one cycle to allow up to four cycles of transfer.
- d. Arbiter does not assert another grant until the bus is idle.

**Step 2:** Graphically illustrate the interface behavior.

**Figure 3-2 Request/Grant/Transfer Relationship for Example 2**



**Step 3:** Using our operational description and graphical waveform, define the detailed temporal requirements for the interface.

- a. Grant should not be asserted until the bus is idle (transfer is done).
- b. Only one grant can be issued at one time.
- c. Grant should be asserted if there is any request pending and the bus is idle.
- d. Grant should not be asserted if there is no request for that port (need to keep track of the number of pending requests).

### Example 3. Latency Considerations for Request and Grant Signals

Although most arbiter interface requirements are relatively simple, there is often latency between the requests at the input to the actual arbiter. Similarly, there is usually some latency after the grant is issued before it is visible at the interface. To model and verify the proper behavior, we must register the request and/or grants to take into account the latency. This is especially apparent in the following example.

## Weighted Priority Arbiter Implementation Interface


---

**Step 1:** Generally describe the operational sequence of events for the interface.

- a. Request is asserted until grant is provided. Otherwise, if request is de-asserted before the grant, it means the client is abandoning the request.
- b. Requests take three cycles to propagate to the internal arbiter, hence there is a minimum three-cycle latency before the request is processed.
- c. Arbiter asserts grant for one cycle.
- d. Transfer signal to indicate bus is busy has one-cycle latency to the arbiter.

**Step 2:** Graphically illustrate the interface behavior.

**Figure 3-3 Request with Latency of Three**



**Step 3:** Using our operational description and graphical waveform, define the detailed temporal requirements for the interface.

The requirements are similar to the ones described in the previous example. However, there are latency indications. Let us consider the following requirements.

- a. Grant should not be asserted without a request.
- b. Grant should be given to a request within three cycles if the bus is idle.

Consider requirement 1: *Grant should not be asserted without a request.* According to [Figure 3-3](#) on page 15, the earliest grant can be asserted is three cycles after request. Asserting the grant earlier contributes to a violation, even though the request is asserted due to internal latency. To properly capture the requirement, we should register the request three times (denoted by `req_d1`, `req_d2`, and `req_d3`) and use `req_d3` as the actual request that is supposed to trigger a grant. If a grant is asserted even one cycle earlier, it is a violation because the internal arbiter will not have seen the request yet. We model this behavior as:


## Weighted Priority Arbiter Implementation Interface

---

```
wire err_gnt_wo_req = gnt[0] & ~req_d3[0];
```

Now we consider requirement 2: *Grant should be given to a request within three cycles if the bus is idle.* The waveform shown in [Figure 3-4](#) on page 16 includes the transfer signal that indicates the bus is busy.

**Figure 3-4 Request with Latency of Three and Transfer with Latency of One**



Here the signal `xfer` is de-asserted (indicates that the bus is idle) after the first cycle. However, since there is a one-cycle latency between transfer and grant, we must use the registered version of `xfer` to verify the grant condition. It is also possible that even though the bus is idle, `gnt[0]` is not asserted because a different port is getting the grant. Therefore, the requirement must include conditions that no other grant is asserted. The following is a sample requirement for a 16-port arbiter.

```
wire err_gnt_not_assert = ~xfer_d1[0] & req_d3[0] & gnt[15:1]==15'b0 & ~gnt[0];
```



---

## Testplan

---

This section presents the high-level requirements on your weighted priority arbiter. A formal testplan includes a high-level requirements checklist, which is the list of requirements you plan to formally verify. Requirements tables list the specific signal names coded in the high-level requirements model and a description of what the requirements check.

Refer to the following tables:

- [Requirements Set 1 – General Arbiter Requirements](#)
- [Requirements Set 2 - Weighted Priority Requirements](#)

**Note:** You will create an additional table that lists interface requirements (see [Table 4-3, “Requirements Set 3 – Arbiter Interface Requirements”](#)).

**Table 4-1 Requirements Set 1 – General Arbiter Requirements**

Requirement Name	Summary
err_req_no_gnt	When a client generates a request, and no other client has generated a request, then the next grant must be issued to the requesting agent.
err_gnt_wo_req	No grant is issued without a request.
err_multiple_gnt	No more than one grant can be issued at one time.

### Weighted Priority Requirements

Since the requirements for a weighted priority arbiter are heavily dependent on a specific implementation (that is, the intended throughput), it is not possible to give a generic English description for the requirements. However, in this section, we give the specific English requirements for the weighted priority arbiter example we introduced in the Overview (see [“Overview”](#) on page 7). You should be able to take the concepts we present in this section and

## Weighted Priority Arbiter Implementation Testplan

modify them to meet your specific requirements for your own implementation of a weighted priority arbiter.

**Table 4-2 Requirements Set 2 - Weighted Priority Requirements**


### Requirements

When `Port0` to `Port3` issue a request, they should get a grant before getting 10 other transfers.

When `Port4` to `Port15` issue a request, they should get a grant before getting 20 other transfers.

The number of transfers (grants) for a weighted priority requirements scheme can be generalized to  $\text{total\_weight}/\text{weight\_of\_port}$ . In this example, the total weight is 20. The weight of `Port1` is two (two locations out of the 20 total weight) and `Port4` is one.

**Figure 4-1 Round-Robin Arbiter Example**



**Note:** If your arbiter has multiple arbitration schemes, simply combine the requirements from the different arbiter formal testplan sections.

## Weighted Priority Arbiter Implementation

### Testplan

---

#### Table 4-3 Requirements Set 3 – Arbiter Interface Requirements

---

This list includes temporal interface requirements for your specific arbiter. Since this list is very specific to your design, you will create it. Refer to the Interface section (see [“Interface”](#) on page 11) for examples of various request/grant interface requirements.

---

# Weighted Priority Arbiter Implementation

## Testplan

---

---

## Requirements

---

In this section we use the interface signals (see [“Interface Names”](#) on page 12) and the set of requirements identified in the Testplan (see [Table 4-2](#) on page 18) to create the high-level requirements models that monitor design block input and output signals.

This section contains the Verilog® source code for a weighted priority arbiter requirements model. This high-level requirements model specifies the behavior for a weighted priority arbiter, where the “weight” of each port is evenly distributed. For more generic requirements, refer to the application note “Dynamic Priority Arbiter Implementation” and consult the Requirements section.

```

/* ----- */
/* - JASPER_weight_arb - */
/* - - */
/* ----- */

`define NUM_PORT 8

module weight_arb (clk, rstN, req, gnt, p1, req2gnt_latency,
                  total_weight, weight_0, weight_1, weight_2, weight_3,
                  weight_4, weight_5, weight_6, weight_7);

    input          clk;
    input          rstN;
    input [ `NUM_PORT-1:0 ] req;
    input [ `NUM_PORT-1:0 ] gnt;

    // Select one arbitrary port for verification

    input [2:0]      p1;
    input [2:0]      req2gnt_latency;

    // The latency from request asserted to grant asserted

```

## Weighted Priority Arbiter Implementation Requirements

---

```
input [7:0]          total_weight;
input [2:0]          weight_0, weight_1, weight_2, weight_3,
                    weight_4, weight_5, weight_6, weight_7;

//  The weight factor for each port, assuming
//  only values 0, 1, 2, and 4 are allowed

reg               req_p1;
reg               req_d1_p1, req_d2_p1, req_d3_p1, req_d4_p1, req_d5_p1,
                    req_d6_p1, req_d7_p1;
reg [2:0]         weight_p1; //the weight of port p1

//  Maximum number of grants before p1 should get a grant

wire [7:0] max_p1_gnt_delay =
    (weight_p1==3'b1) ? (total_weight-1'b1) :
    (weight_p1==3'd2  ? ({1'b0, total_weight[7:1]}) :
    (weight_p1==3'd4  ? ({2'b0, total_weight[7:2]}) : 8'b0));

reg [7:0]  num_gnt;
always @(posedge clk) begin
    if (rstN)
        num_gnt <= 5'b0;
    else if (gnt[p1])
        num_gnt <= 5'b0;
    else if (req_p1 & gnt!=`NUM_PORT'b0 & ~gnt[p1])
        num_gnt <= num_gnt + 1'b1;
end

//  Selecting the request that corresponds to the grant

always @(req_d1_p1 or req_d2_p1 or req_d3_p1 or req_d4_p1 or
        req_d5_p1 or req_d6_p1 or req_d7_p1 or req or p1 or
        req2gnt_latency) begin

    case (req2gnt_latency)
        0: req_p1 = req[p1];
        1: req_p1 = req_d1_p1;
        2: req_p1 = req_d2_p1;
        3: req_p1 = req_d3_p1;
```

## Weighted Priority Arbiter Implementation Requirements

---

```
        4: req_p1 = req_d4_p1;
        5: req_p1 = req_d5_p1;
        6: req_p1 = req_d6_p1;
        7: req_p1 = req_d7_p1;
    endcase

end

// Registering request to reflect the proper request to
// grant latency

always @(posedge clk) begin

    if (!rstN) begin
        req_d1_p1 <= 1'b0;
        req_d2_p1 <= 1'b0;
        req_d3_p1 <= 1'b0;
        req_d4_p1 <= 1'b0;
        req_d5_p1 <= 1'b0;
        req_d6_p1 <= 1'b0;
        req_d7_p1 <= 1'b0;
    end
    else begin
        req_d1_p1 <= req[p1];
        req_d2_p1 <= req_d1_p1;
        req_d3_p1 <= req_d2_p1;
        req_d4_p1 <= req_d3_p1;
        req_d5_p1 <= req_d4_p1;
        req_d6_p1 <= req_d5_p1;
        req_d7_p1 <= req_d6_p1;
    end

end

always @(weight_0 or weight_1 or weight_2 or weight_3 or
        weight_4 or weight_5 or weight_6 or weight_7 or p1)
begin

    case (p1)
        0: weight_p1 = weight_0;
        1: weight_p1 = weight_1;
```

## Weighted Priority Arbiter Implementation Requirements

---

```
        2: weight_p1 = weight_2;
        3: weight_p1 = weight_3;
        4: weight_p1 = weight_4;
        5: weight_p1 = weight_5;
        6: weight_p1 = weight_6;
        7: weight_p1 = weight_7;
    endcase

end

//  Grant should not be issued to a port with 0 weight

wire err_weight_1 = gnt[p1] & weight_p1==3'b0;

//  The maximum latency for grant should not exceed the
//  total-weight/weight_p1

wire err_weight_2 = req_p1 & gnt != `NUM_PORT'b0 &
    (weight_p1==3'b0) & (num_gnt==max_p1_gnt_delay);

endmodule
```



---

## Strategy

---

This section shows techniques and strategies that address potential arbiter performance issues. Several advanced techniques are generally used to formally verify arbiters. We recommend you review all the suggestions in this section. However, it is not critical to understand all the details of the various techniques prior to starting your proof. You may choose to apply a few of our suggestions to proactively prevent performance issues during your proof. Alternatively, you might choose to wait and apply a few of our suggestions later only if you are unable to complete your proof due to a performance problem.

This section includes the following topics:

- [Verification Plan](#)
- [What Ports Should You Monitor as Part of Your Requirements?](#)
- [Speeding Up Your Proof](#)

## Verification Plan

The Verification Plan contains a set of optional “Restriction Definitions” and the recommended “Verification Steps.” The combination of restrictions and steps forms an effective methodology.

### Restriction Definitions

For an immature design, you might need restrictions. As the design matures, you might choose to reduce the restrictions or skip them altogether. Restrictions allow you to verify basic functionality and mainstream problems before complicated corner-case bugs. They also allow you to verify a design before it is complete.

For example, a designer is working on a 64-port arbiter with round-robin arbitration and three levels of strict priority. To simplify, you might implement a one-priority, round-robin arbiter with two ports before expanding to more ports and more fixed priorities. Therefore, it is probably beneficial to first set the restrictions to arbitrate only between two ports.

## Weighted Priority Arbiter Implementation Strategy

---

1. Only one port has an active request.
2. Only two ports have active grants.
3. Only one arbitration scheme is active at a time (for arbiters containing multiple arbitration schemes).

### Verification Steps

The following lists recommended steps for your proof.

1. Prove [Requirements Set 1](#) with [Restriction Definition 1](#), [Restriction Definition 2](#), and [Restriction Definition 3](#).
2. Prove [Requirements Set 3](#) with [Restriction Definition 1](#), [Restriction Definition 2](#), and [Restriction Definition 3](#).
3. Prove [Requirements Set 1](#) with [Restriction Definition 2](#) and [Restriction Definition 3](#).
4. Prove [Requirements Set 3](#) with [Restriction Definition 2](#) and [Restriction Definition 3](#).
5. Prove [Requirements Set 2](#) with [Restriction Definition 2](#) and [Restriction Definition 3](#).
6. Repeat step 5 for all different arbitration schemes if applicable.
7. Prove [Requirements Set 1](#) with [Restriction Definition 3](#).
8. Prove [Requirements Set 3](#) with [Restriction Definition 3](#).
9. Prove [Requirements Set 2](#) with [Restriction Definition 3](#) (for all arbitration schemes if applicable).
10. Prove [Requirements Set 1](#) with no restrictions.
11. Prove [Requirements Set 3](#) with no restrictions.
12. Prove [Requirements Set 2](#) with no restrictions.

### What Ports Should You Monitor as Part of Your Requirements?

When creating an arbiter requirements model, we often take advantage of the following property:

## Weighted Priority Arbiter Implementation Strategy

---

*If each individual port of an arbiter does not violate any of its requirements, then an arbiter that consists of multiple ports will never violate any of the requirements.*

That is, for many arbiter requirements, it is sufficient to prove the requirement with respect to a single port separately versus attempting to prove the requirement with respect to all the ports at once. Furthermore, writing a requirement for a single port is usually easier than writing a requirement that attempts to capture the behavior of all the ports.

Consider the following example:

### ■ Specification

A grant should not be asserted for any particular port when there was not a corresponding request issued for that same port.

### ■ Requirements Model

```
wire err_gnt_wo_req_1 = (gnt != `NUM_PORT'b0) & ((gnt & req) == `NUM_PORT'b0);
```

In the previous example, we are modeling the required behavior using a Boolean expression, which is monitoring the occurrence of grant (that is, the n-bit `gnt` variable is non-zero). When a grant occurs, the expression checks to see if a corresponding request is missing (that is, the n-bit `req` variable is zero). We model this behavior by "AND-ing" n-bit `req` (request) and `gnt` (grant) variables while making the assumption that grant will be one-hot (which you should prove as a separate requirement).

As you can see, trying to specify all ports at once makes a simple requirement more complicated than necessary. Furthermore, specifying all ports as part of the requirement can increase the proof run time since exploring a larger state space might be required to complete the proof.

Now consider the following alternative:

```
wire err_gnt_wo_req_2 = gnt[p1] & ~req[p1];
```

where the `p1` variable (which is part of our requirement model) is used to select any arbitrary arbiter port from the n-bit `gnt` and `req` signals. You can see that the new requirement is simpler and straight forward. To ensure that all ports are proven correctly, we recommend that you use one of the two following techniques.

### ■ Technique 1: Use the `assume -constant` command.

```
assume -constant p1
```

## Weighted Priority Arbiter Implementation Strategy

---

The `assume -constant` command ensures that the port number is held to an arbitrary constant value throughout the evaluation of the proof. That is, all arbitrary values are explored, but for any particular single trace (path) the value will not change, which means that the formal tool will not return a false failure due to changing a port selection for a request or grant in the middle of its evaluation.

The advantage of using this constraint to specify an arbitrary value is that we are able to verify our requirement for all ports without the need to verify each port separately (since JasperGold® Apps will try all possible combinations). Alternatively, you can use Technique 2 to set `p1` to a constant and go through all ports sequentially.

### ■ **Technique 2:** Create a script to prove each port separately.

```
# Prove port 0
assume -name assume_port_0 p1==3'd0
prove ~err_gnt_wo_req_1
assume -remove assume_port_0
#
# Prove port 1
assume -name assume_port_1 p==3'd1
prove ~err_gnt_wo_req_1
assume -remove assume_port_1
.
.
.
```

The advantage of the second technique is that each proof is simpler compared to the first technique's `assume -constant` command. However, it comes at a cost of less automation since the second technique requires you to manually create the script to check each port. Therefore, we recommend that you initially try Technique 1.

**Note:** When proving arbiter requirements, it is often easier to prove an arbitrary pair of ports that is specified by two variables `p1` and `p2` (similar to Technique 1) than attempt to prove all ports as part of a single requirement.

## Speeding Up Your Proof

For a general discussion on techniques to speed up your proof, refer to the [“Appendix: Techniques and Strategies”](#) on page 33. To go directly to a specific topic of interest, click on one of the following cross-reference links:

### ■ [Taking Advantage of Design Symmetry](#) on page 34

## Weighted Priority Arbiter Implementation Strategy

---

- [Multiplexer Abstraction](#) on page 35
- [FIFO Abstraction](#) on page 37
- [Induction](#) on page 40
- [Counter Abstraction](#) on page 41
- [Separating Decoding Logic](#) on page 43

## Weighted Priority Arbiter Implementation Strategy

---

---

## Coverage

---

In this section, we define a set of coverage goals you can use to confirm that your coverage objectives were met during the proof.

After your requirements prove true, you should always perform a sanity trace and cover on all major functionality. It is important to:

- Ensure that you are proving what you think you are proving
- Cover the function and exit from that function
- Ensure that specific states or conditions in your high-level requirements model have not been over-constrained by any assumptions

For example, we can get the coverage on a request being asserted. However, the proper coverage includes a request being asserted *and* de-asserted. Similarly, we should check the coverage of the priority going from one level to another, rather than just one priority level at one time. To support these objectives, this section lists coverage goals you will check after a requirement proves true.

**Note:** Refer to the [command reference](#) for details on the `visualize`, `cover`, and `prove` commands.

### Coverage Set 1

- Coverage on all requests, grants, and transfers (if applicable).
- Coverage on all priority levels, all credit/weight levels.

### Coverage Set 2

- Coverage on all dynamic allocation processes: refresh, out-of-credit, and so forth.

## **Weighted Priority Arbiter Implementation Coverage**

---



---

# Appendix: Techniques and Strategies

---

In some cases, your proof can benefit from an advanced methodology technique to overcome complexity problems during a proof. For example, various abstraction techniques are required when the number of arbiter clients is high. Abstraction is one technique that can simplify many complex structures without compromising the integrity of the proof. In this appendix, we introduce a few advanced techniques that can be useful for speeding up your proof. We suggest you review all the suggestions in this appendix; however, it is not critical to understand all the details of the various techniques prior to starting your proof. You might choose to apply a few of our suggestions to proactively prevent performance issues during your proof. Or you might choose to wait and apply our suggestions later if you are unable to complete your proof due to a performance problem.

This appendix includes the following techniques for speeding up your proofs:

- [Taking Advantage of Design Symmetry](#)
- [Multiplexer Abstraction](#)
- [FIFO Abstraction](#)
- [Induction](#)
- [Counter Abstraction](#)
- [Separating Decoding Logic](#)

## Taking Advantage of Design Symmetry

Many types of designs (such as pointer-based, round-robin arbiters) have a symmetrical implementation. That is, the logic associated with each port has the exact same RTL structure. Generally, the only exception to a symmetrical arbiter implementation is the initial priority value for each port (that is, you will give one port the highest priority).

Symmetry allows us to efficiently prove requirements on a pair of symmetrical paths (for example, ports of an arbiter) while giving us confidence that the requirement is true on all ports due to the symmetrical implementation. If you are confident about the symmetry among ports in the implementation of your design, then we recommend you take advantage of this symmetry as part of your proof.

### Arbiter Example

Consider the following Verilog® code fragment for an arbiter that has a symmetrical implementation per port with the exception of the initial priority value:

```
// pointer to determine round-robin
reg [3:0] priority_pointer;
// priority port

reg [7:0] grant;

always @(posedge clk) begin
    if (reset) begin
        priority_pointer <= 4'd0;
        grant <= 8'd0;
    end
    else if (grant != 8'd0) grant <= 8'd0;
    else if (request[priority_pointer]) begin
        priority_pointer <= priority_pointer + 5'd1;
        grant[priority_pointer] <= 1'b1;
    end
    else if (request[priority_pointer+5'd1]) begin
        priority_pointer <= priority_pointer + 5'd2;
        grant[priority_pointer + 5'd1] <= 1'b1;
    end
    else if (request[priority_pointer+5'd2]) begin
        priority_pointer <= priority_pointer + 5'd3;
```

## Weighted Priority Arbiter Implementation

### Appendix: Techniques and Strategies

---

```
    grant[priority_pointer + 5'd2] <= 1'b1;
end
.
.
.
end
```

In the previous example, you can see that the logic associated with each port is symmetrical since each n-bit variable assignment follows the same RTL structural path. The only part of this implementation that is not symmetrical is the reset value of `priority_pointer`, which points to `port0` after reset (that is, it has the highest priority after reset). To allow true symmetry for our proof, we can use the `abstract -init_value` command to ignore the initial value.

```
abstract -init_value priority_pointer
```

This command permits all possible initialization values for the signal `priority_pointer` to be explored, which restores the true symmetry to the circuit. In the induction section (below), you will see that this command is useful for proving other aspects of an arbiter.


## Multiplexer Abstraction

Multiplexers are frequently used to implement arbitration schemes that involve multiple ports. For the case where the port count is high, multiplexers increase the complexity of the proof and present a problem. However, most arbiter requirements can be represented in terms of `p1` and/or `p2` instead of specifying all ports or a fixed port number directly. Refer to any of the arbiter application notes (for example, “Fairness Arbiter Implementations”) to learn about a technique for specifying a few arbiter requirements by using a fixed (but arbitrary) port `p1` as part of your specification.

Since our objective is to identify a bug in the control logic for our arbiter, we can take advantage of the technique we present in this section to reduce the complexity of the multiplexer requirements model (as well as the design itself during the proof).

Consider the following example:

**Figure A-1 Multiplexer Abstraction**



In [Figure A-1](#) on page 36, we illustrate the concept of reducing the complexity of the multiplexer. For this example, we have taken an N-to-1 multiplexer (shown on the left) and reduced it to a 3-to-1 multiplexer (shown on the right). Our abstraction is safe since, for our proof, we only need to distinguish two arbitrary ports  $p_1$  and  $p_2$  (represented as “Port P1” and “Port P2” in the figure). Since we are not distinguishing the behavior of all other ports as part of our requirement, we can represent the behavior of all other ports using a single port we have labeled “Port Other” (refer to [Figure A-1](#) on page 36). JasperGold Apps explores the behavior of the ports we must distinguish (that is,  $p_1$  and  $p_2$ ) in the context of the behavior for all the other ports that we do not need to distinguish.

For example, consider the original model:

```
always @(port[0] or port[1] or port[2] or port[3] or port[4] or port[5] or
port[6] or port[7] or round_robin_ptr) begin
  case (round_robin_ptr)
    0: port_out = port[0];
    1: port_out = port[1];
    .
    .
    .
    15: port_out = port[15];
  endcase
end
```

## Weighted Priority Arbiter Implementation

### Appendix: Techniques and Strategies

---

There are two abstraction techniques we can use to simplify this code when proving the requirements in the context of two arbitrary ports, `p1` and `p2`.

#### ■ Technique 1: RTL Code Version of Abstracted Multiplier Model

```
wire [2:0] port_other;

always @(port[0] or port[1] or port[2] or port[3] or port[4] or port[5] or
port[6] or port[7] or round_robin_ptr or p1 or p2) begin
    if (round_robin_ptr==p1) port_out=port[p1];
    else if (round_robin_ptr==p2) port_out=port[p2];
    else port_out = port_other;
end
```

Using the above model, you should add the following Tcl commands to your script:

```
assume {port_other!=port[p1]}
assume {port_other!=port[p2]}
```

#### ■ Technique 2: Tcl Script Version of Abstracted Multiplexer Model

```
stopat {port_out}
assume {round_robin==p1 => port_out==port[p1]}
assume {round_robin!=p1 => port_out!=port[p1]}
assume {round_robin==p2 => port_out==port[p2]}
assume {round_robin!=p2 => port_out!=port[p2]}
```

## FIFO Abstraction

Although in general, FIFOs are more likely to be found in the datapath portion of a design rather than the actual arbiter, there are some arbiter implementations that take advantage of a FIFO structure. Consider a round-robin arbiter that uses registers to store the priority for each port. Every time a grant is provided, the priority changes, which affects at least one of the priority registers, and potentially affects all of the priority registers. This priority scheme using multiple registers, in a way, behaves similar to a FIFO (except it may not always behave in a first-in first-out fashion). Since the set of priority registers can be large, we might need to replace this set of large registers with an abstracted model. The following example demonstrates this technique.

## Example

Consider a simple 32-location, circular FIFO-type structure that is keeping track of some type of information (for example, its priority) related to each port of our arbiter. Information enters our FIFO at location 31 and is shifted as new information arrives until it exits at location 0. Essentially, for an arbiter design, the information indicating high priority at location 0 will move to location 31 using this circular FIFO structure. By using the basic ideas related to `p1` and `p2`, that is, if each individual port of an arbiter does not violate any of its requirements, then an arbiter that consists of multiple ports will never violate any of the requirements, we are able to dramatically simplify the FIFO and reduce the complexity of the design for our proof.<sup>1</sup> Instead of keeping track of the content for every FIFO location, we only need to keep track of the information for `p1` and `p2`, along with their locations within the FIFO. You will see that this reduces our large 32-location FIFO down to four smaller registers and enables us to prove complex properties related to the FIFO controller.

The following Verilog RTL code fragment demonstrates the original RTL code for the FIFO.

```
reg [15:0] fifo0;
reg [15:0] fifo1;
reg [15:0] fifo2;
.
.
.
reg [15:0] fifo31;

always @ (posedge clk) begin
  if (reset) begin
    fifo0 <= 16'h0;
    fifo1 <= 16'h0;
    .
    .
    .
    fifo31 <= 16'h0;
  end
  else if (fifo_en) begin
    fifo31 <= new_info;
    fifo30 <= fifo31;
    fifo29 <= fifo30;
```

1. Refer to [“What Ports Should You Monitor as Part of Your Requirements?”](#) on page 26 for a discussion of this concept.

## Weighted Priority Arbiter Implementation

### Appendix: Techniques and Strategies

---

```
.
.
.
fifo0 <= fifo1;
info_out <= fifo0;
end
end
```

The following Verilog RTL code fragment demonstrates our FIFO abstraction technique.

```
reg [15:0] p1_info;
reg [15:0] p2_info;
reg [4:0]   p1_loc;
reg [4:0]   p2_loc;

always @(posedge clk) begin
  if (reset) begin
    p1_info <= 16'h0;
    p2_info <= 16'h0;

    //initial position of p1 within the fifo

    p1_loc <= p1_init;

    //initial position of p2 within the fifo

    p2_loc <= p2_init;
  end
  else if (fifo_en) begin
    p1_loc <= p1_loc - 5'h1;
    p2_loc <= p2_loc + 5'h1;
    if (p1_loc <= 5'h0) begin
      p1_info <= new_info;
      info_out <= p1_info;
    end
    else if (p2_loc <= 5'h0) begin
      p2_info <= new_info;
      info_out <= p2_info;
    end
  end
end
end
```

## Induction

One of the main advantages of using JasperGold Apps is that it provides the trace that demonstrates the requirement violations with a minimum number of clock cycles. This feature makes the full verification a lot faster and a lot easier to debug. However, certain structures can unnecessarily increase the number of cycles required to demonstrate the bug. These situations are usually caused by structures where the function being evaluated follows this type of relationship:

$$f(x+1) = g[f(x)]$$

where  $f[0]$  is constant. One of the most common examples is a counter. Refer to the “Overview” section of the application note “Fair Arbiter Implementations” to see examples of functions that follow this relationship.

For example, the logic that generates the priority pointer values for a round-robin arbiter follows this relationship. Similarly, in the FIFO abstraction example shown above, the logic that generates location pointer values for  $p1$  and  $p2$  follows this relationship. If we can ignore the initial value for these pointer examples, thus allowing all possible initialization conditions, we can shorten the number of cycles required to verify the requirement by using induction.

## Example

The round-robin pointers are initialized to zero. To prove the requirements, the design must go through the states containing every pointer value. If the pointer goes from 0 to 31, it will take at least 32 transactions to complete the transfer. However, if we free up the initial value of the pointers to allow any initial values, it only takes one transaction to reach all pointer values from 0 to 31. Consequently, the number of cycles needed to prove the requirements is significantly smaller.

Being able to prove the requirements without needing the initial value also implies that the design is symmetrical. That is, if the design is truly symmetrical, we only need to prove the interface requirements by monitoring one port and fairness requirements by monitoring one pair of ports.

The next section demonstrates how to use induction to simplify a proof that involves a large counter.



## Counter Abstraction

Counters are one of the most common components used in a design. If a counter is large, it can be a challenge for any formal verification tool because it can cause the number of evaluation iterations to be very high.

In an arbiter, there are several places that require counters. Dynamic arbiters, for example, typically use a counter to keep track of the arbitration change events. Fortunately, it is usually possible to abstract these counters to reduce the reachable states without compromising the proof.

There are two main counter abstractions: *Counter Induction* and *Counter Reduction*.

### Counter Induction

Counter induction uses *induction* to abstract the counter. For example, consider a dynamic priority arbiter that requires us to prove that the credit/weight is incremented, reset, and decremented correctly. If we follow the actual logic where the credit is reset to a fixed value, we will need to go through many cycles to explore all possible credit values. Alternatively, we can partition the requirement into two parts:

- The credit register is initialized correctly.
- The credit register is incremented and decremented correctly.

The first requirement can be modeled as follows:

```
// "credit_init" is the expected initialization value of credit and
// "prev_rst" is the registered version of reset, hence indicate the value
// of credit during the very first cycle
```

```
wire err_credit_reset = prev_rst & (credit!=credit_init);
```

The second requirement can be modeled as follows:

```
wire credit_increment = credit_inc & (credit!=prev_credit+1);
wire credit_decrement = credit_dec & (credit!=prev_credit-1);
```

In addition, for the second requirement you must let the initial credit value be free (that is, it is uninitialized and allowed to take on any value). Add the following command to your JasperGold Tcl script to let the initial credit value be free:

```
abstract -init_value {credit}
```

## Counter Reduction

Use the second type of counter abstraction when a specific counter value triggers a specific event (for example, a timeout counter). Since most of the counting values are uninteresting, and are only used as a sequence to generate the next interesting counter value, we can perform an abstraction to reduce the number of states the formal engine needs to evaluate.

For example, assume our design has a counter that calculates the elapsed time before the design must refill its credit for a credit-based arbiter. For this type of design, there are only two counter values that are interesting (that is, zero and the timeout count to generate the credit refill). All other counter values have no influence or impact on the rest of the design. Hence, we can use an abstraction technique that will *fast forward* the counter values that are uninteresting and thus reduce the state space needed to prove the design's high-level requirements.

Original counter model:

```
reg [15:0] count;

always @(posedge clk) begin
    if (rst) count <= 16'h0;
    else if (count == credit_refresh_cnt) count <= 16'h0;
    else count <= count + 16'h1;
end
```

Abstracted counter model:

```
reg abstract_cnt;
reg [15:0] nxt_count;

always @(posedge clk) begin
    if (rst) abstract_cnt <= 1'b0;
    else if (count==credit_refresh_cnt) abstract_cnt <= 1'b0;
    else abstract_cnt <= 1'b1;
end
```

For the abstracted counter, add the following to your JasperGold Tcl script:

```
stopat count
```

## Weighted Priority Arbiter Implementation

### Appendix: Techniques and Strategies

---

```
# count can assume any value less than or equal to credit_refresh_cnt

assume {abstract_cnt => (count<=credit_refresh_cnt)}

# count will be set back to zero

assume {~abstract_cnt => (count == 16'h0)}
```

In our previous example, during reset, the value of `abstract_cnt` is set to zero, which in turn sets `count` to zero with the Tcl assumption. After that, `abstract_cnt` is set to one, which allows `count` to assume any (and all) values between zero and the `credit_refresh_cnt` with the Tcl assumption. If there are no meaningful states for JasperGold Apps to analyze between zero and `credit_refresh_cnt`, then `count` will not arbitrarily assume a value that is less than `credit_refresh_cnt` for any time longer than necessary. Hence, `count` will assume `credit_refresh_cnt` (an interesting state) sooner without having to sequence through all the counter values. When `count` assumes the value of `credit_refresh_cnt`, `abstract_cnt` resets back to zero, which in turn resets `count` back to zero.

One advantage for this type of counter abstraction is that if we miss some meaningful counts (counts that impact the rest of the design) when we are creating our abstract model, our proof returns a false with a counterexample. Thus, this form of abstraction might give you a false negative, but it does not give you a false positive. If there are more counter values needed in our abstracted model, we can then declare more states and map these states onto states in our abstracted model. This type of counter abstraction is also useful when counting the number of remaining credits in the credit-based dynamic arbiter.

## Separating Decoding Logic

Separating the decoding logic from the logic used to ensure fairness is especially useful for arbiters that involve dynamic prioritizations and fixed priority arbiters. The essential idea is that for many arbiters there is decoding logic that is responsible for generating the various priority levels based on multiple factors.

For example, assume we are designing an arbiter where the priority of a specific port is increased from normal to high whenever the input FIFO for that port reaches the high watermark. The priority increases further from high to highest when the FIFO is full. All ports with the same priority levels (normal, high, and highest) are arbitrated using a round-robin scheme. However, each priority level has strict priority over the lower priority levels.

To write a high-level requirement for this combined round-robin and priority scheme can be complicated. In addition, proving this combined scheme can be complicated for a formal tool. However, if we separate the verification problem into two parts, we can simplify the coding of

## Weighted Priority Arbiter Implementation

### Appendix: Techniques and Strategies

---

our high-level requirement and simplify the proof. The requirement can be partitioned as follows:

- Verify that the arbitration requirement is met.
- Verify that the generation of the priority is correct.

For example, to separate the logic that generates the priority from the logic that implements the arbitration scheme, add the following command to your JasperGold Tcl script when you prove that the arbitration requirement is correct:

```
stopat {priority[15:0]}
```

Then, in your requirements model, code the following abstract priority model:

```
// If there is at least one port with high priority, the effective
// requests are the ports with request and high priority.
// Otherwise, it is the same as a request
// if there is no high priority request.

reg [15:0] priority_req = (priority==16'b0) ? (priority&req) : req;

// Use priority_req in place of regular request for any
// requirements that are applicable
```

There are also a few assumptions that might be needed to model the behavior of the abstracted priority model correctly. For example, whenever a priority is high, then the priority should not change to low priority prior to receiving a grant. This assumption requires additional modeling in the requirement as shown:

```
// registered previous value of priority

reg [15:0] prev_priority;

// registered previous value of grant

reg [15:0] prev_gnt;

always @(posedge clk) begin
  if (rst) begin
    prev_priority <= `NUM_PORT'h0;
    prev_gnt <= `NUM_PORT'h0;
  end
end
```

## Weighted Priority Arbiter Implementation

### Appendix: Techniques and Strategies

---

```
else begin
    prev_priority <= priority;
    prev_gnt <= gnt;
end
end
```

Then add the following assumptions to your Tcl script:

```
assume {(~prev_gnt[0] & prev_priority[0]) => (priority[0]==1'b1)}
assume {(~prev_gnt[1] & prev_priority[1]) => (priority[1]==1'b1)}
.
.
.
```

# **Weighted Priority Arbiter Implementation**

## Appendix: Techniques and Strategies

---